AgentGuard

Cybersecurity for the Agentic Era

The runtime security platform that monitors, constrains, and audits AI agents — before they go catastrophic.

Confidential — February 2026
Business Case v2.0 — Updated with Market Intelligence

The Market Is Exploding

$6.34B	3×	$54M	$30B+
AI Security Funding in 2025	YoY Funding Growth	Avg Deal Size (up from $34M)	Projected TAM by 2028

AI agent security is the fastest-growing segment in cybersecurity. In 2025 alone:

- **7AI** raised $130M Series A at $700M valuation
- **Noma Security** closed $100M Series B
- **Sweet Security** raised $75M Series B
- **WitnessAI** raised $58M specifically for agentic AI governance
- **Snyk acquired Invariant Labs** — the clearest signal that agent security is a real category

> **The signal is unmistakable:** The biggest names in developer security are racing to own agent security. The window for category creation is 12–18 months.

The Problem

AI agents are moving from demos to production. They book flights, write code, manage infrastructure, handle finances, and make decisions with real-world consequences.

But there's no security layer between "agent was given a task" and "agent did something catastrophic."

The Full Threat Landscape

Threat	Description	Severity
Prompt Injection	Agents manipulated by malicious inputs to act against owner's intent	CRITICAL
Data Poisoning via Memory	Persistent backdoors in RAG stores influencing all future decisions	CRITICAL
Cross-Agent Injection	Compromised agent injects malicious instructions through the agent graph	CRITICAL
Tool Confusion	Agent tricked into calling wrong tool or passing malicious parameters	HIGH

Goal Drift	Subtle deviation from instructions over long reasoning chains	HIGH
Privilege Escalation	Agents acquiring access beyond what was intended	HIGH
Multi-Agent Cascade	Agents "fixing" each other's fixes in emergent failure loops	HIGH
Shadow Actions	Side-effects the user never asked for — exfiltration, spending, API calls	HIGH

> **What a CISO says:** *"I have 200 agents deployed across 6 departments. I can't tell you what they accessed yesterday, whether any violated our data handling policies, or if one is currently exfiltrating PII. My board asks about AI risk every quarter and I have nothing to show them."*

Competitive Landscape — Deep Analysis

Despite $6.34B flowing into AI security, **critical gaps remain unfilled**. Here's who's playing and what they miss:

Player	Focus	Funding	The Gap
Lakera 🇨🇭	Prompt injection defence, data leakage	Series A+	No agent lifecycle governance, no kill switch, no multi-agent coordination
Protect AI 🇺🇸	Model scanning, red teaming, runtime monitoring	Series B	ML model security focus. Less on agentic behaviour, tool call policies
CalypsoAI 🇺🇸	Red-team, defend, observe — full lifecycle	Series B	Closest competitor. Less on declarative policy engine, HITL gates, multi-agent
Invariant Labs	Agent behaviour inspection, guardrails, MCP security	Acquired by Snyk	Now part of Snyk ecosystem — validates category but limits independence
Pillar Security 🇮🇱	AI asset discovery, red team, runtime guardrails	Series A	Broad AI security. Agent features exist but aren't primary focus
WitnessAI 🇺🇸	Agentic AI governance	$58M	Governance focus — less on runtime enforcement and incident response
7AI	AI identity & access security	$130M	Identity layer only — no action-level policy, no audit trail, no compliance
Noma Security	AI supply chain & data security	$100M	Supply chain focus — not runtime agent behaviour monitoring
Traditional SIEM	Infrastructure & endpoint monitoring	N/A	Don't understand agent intent, reasoning chains, or agent-specific threats

The 7 Unfilled Gaps — AgentGuard's Opportunity

#	Gap	Why No One Fills It	AgentGuard's Play
1	**Multi-Agent Governance**	Everyone focuses on single agents. Multi-agent orchestration is the 2027+ wave	Cross-agent visibility, collusion detection, hierarchical permissions from day 1
2	**Declarative Policy Engine**	Competitors detect but can't enforce. No "OPA/Rego for AI	YAML/JSON policies, version-controlled, action-level

		agents" exists	enforcement
3	**MCP Security**	Model Context Protocol becoming standard but no standalone MCP security solution	MCP-native security layer — inspect, validate, and control tool calls via MCP
4	**Agent Zero-Trust**	No one has built mTLS between agents, JIT privilege elevation, continuous verification	Full zero-trust architecture: agent identity, microsegmentation, least privilege
5	**Cross-Framework Compatibility**	Most tools are model/framework-specific	Framework-agnostic middleware — works with any agent runtime
6	**EU AI Act Compliance Automation**	Deadline Aug 2026 — 5 months away. No automated compliance evidence for agentic AI	One-click compliance reports mapping agent behaviour to Articles 9, 12–15
7	**Human-in-the-Loop Gates**	OWASP flagged as top risk. Tools detect but don't intervene	Configurable HITL with approval workflows, escalation, and graceful degradation

Strategic insight: Snyk acquiring Invariant Labs validates the category. But Snyk will integrate it into their developer platform, not build a standalone agent security product. The independent, framework-agnostic position is wide open.

The Solution: AgentGuard

A runtime security platform that monitors, constrains, and audits AI agents — regardless of which model or framework powers them.

🔥 Core Capabilities

1. Policy Engine (the "Firewall")
Declarative YAML/JSON policies. Action-level permissions, budget/rate limits, scope boundaries. Version-controlled. The first "OPA for agents."

2. Real-Time Monitoring (the "SOC")
Every action logged with full context. Anomaly detection, chain-of-thought inspection, risk scoring. SIEM integration (Splunk, Sentinel, Chronicle).

3. Kill Switch & HITL Gates
Instant halt. Human approval for high-risk actions. Configurable fail-closed/fail-open per agent criticality. Graceful degradation.

4. Audit & Compliance Engine
Tamper-evident logging. Pre-built reports: EU AI Act, SOC 2, ISO 42001, DORA, NIS2, HIPAA, PCI-DSS. Court-admissible forensics.

5. Multi-Agent Governance
Cross-agent visibility. Collusion prevention. Hierarchical permissions. Lateral movement detection. Coordination policies.

6. MCP Security Layer
Native Model Context Protocol security. Tool call validation, parameter sanitisation, credential protection. The first MCP firewall.

Architecture — Zero-Trust Agent Security

```
                    ┌──────────────────────────────────────────┐
                    │          AgentGuard Control Plane         │
                    │  ┌──────────┐ ┌──────────┐ ┌──────────┐  │
                    │  │ Policy   │ │ Identity │ │ Threat   │  │
                    │  │ Engine   │ │ & AuthZ  │ │ Intel    │  │
                    │  └──────────┘ └──────────┘ └──────────┘  │
                    │  ┌──────────┐ ┌──────────┐ ┌──────────┐  │
                    │  │ Anomaly  │ │Compliance│ │ MCP      │  │
                    │  │Detection │ │ Engine   │ │ Security │  │
                    │  └──────────┘ └──────────┘ └──────────┘  │
                    └──────────────────────────────────────────┘
                                        │
                                        │
          ┌─────────────────────────────┼─────────────────────────────┐
          │                             │                             │
┌─────────┴─────────┐         ┌─────────┴─────────┐         ┌─────────┴─────────┐
│  Agent A          │         │  Agent B          │         │  Agent C          │
│  (Sandboxed)      │◄──mTLS──►  (Sandboxed)      │◄──mTLS──►  (Sandboxed)      │
└─────────┬─────────┘         └─────────┬─────────┘         └─────────┬─────────┘
          │                             │                             │
┌─────────┴─────────┐         ┌─────────┴─────────┐         ┌─────────┴─────────┐
```

```
    │ Tools (MCP)   │       │ Tools (MCP)   │       │ Tools (MCP)   │
    │ (Proxied)     │       │ (Proxied)     │       │ (Proxied)     │
    └───────────────┘       └───────────────┘       └───────────────┘
```

Business Model

Hybrid Pricing: Platform Fee + Usage + Compliance Modules

Tier	Platform Fee	Usage	Compliance
Free / OSS	$0	3 agents, 10K actions/mo	—
Team	$499/mo	+$0.001/action above 100K	—
Business	$2,500/mo	+$0.0008/action above 500K	$1,500/mo per module
Enterprise	$8K–25K/mo	Custom pricing	$3K–8K/mo bundled
Platform/OEM	Negotiated	Revenue share 12–18%	White-label

Unit Economics

Metric	Team	Business	Enterprise
ACV	$6K	$36K	$120K
Gross Margin	85%	83%	78%
CAC	$2K (PLG)	$18K (inside sales)	$55K (field sales)
LTV (3-yr, 120% NRR)	$20K	$125K	$430K
LTV:CAC	10×	6.9×	7.8×
Payback (months)	4	8	14

Go-To-Market Strategy

Buyer Personas

Persona	Title	Pain Point	Deal Size
Security Champion	CISO / VP Security	"I can't govern what I can't see"	$30K–100K
Platform Builder	VP Eng / Head of AI	"We need guardrails before production"	$10K–40K
Compliance Driver	GRC Lead / DPO	"EU AI Act is 5 months away"	$20K–60K
Innovation Leader	CDO / CTO	"We want to move fast safely"	$15K–50K

Phased Execution

Phase 1 — Developer Love (M1–6)

- Open-source core policy engine (Apache 2.0)
- Integrate top 5 frameworks
- MCP security layer
- KPIs: 5K GitHub stars, 1K weekly OSS users

Phase 2 — Enterprise Wedge (M6–12)

- Compliance-first: EU AI Act deadline messaging
- 25 enterprise design partners
- Zero-trust agent identity, IR playbooks
- KPIs: $500K pipeline, 3–8 paying customers

Phase 3 — Platform Play (M12–24)

- Cloud marketplace listings (AWS, Azure, GCP)
- "AgentGuard Verified" certification
- ASPM dashboard, cross-agent threat correlation
- KPIs: $2M+ ARR, 100+ customers, 140%+ NRR

Financial Projections

Scenario Analysis

Scenario	Prob.	Year 1 ARR	Year 2 ARR	Year 3 ARR
Bull — Category creation, catalytic event	15%	$350K	$4.5M	$18M
Base — Steady execution	55%	$150K	$2.0M	$8.5M
Bear — Adoption stalls, incumbents bundle	30%	$60K	$600K	$2.5M

Base Case Detail

Metric	Year 1	Year 2	Year 3
Enterprise Customers	3	25	100
Team/Business Customers	20	80	280
ARR (exit)	$150K	$2.0M	$8.5M
Headcount	8	18	42
Net Burn	($1.6M)	($2.8M)	($3.9M)
Burn Multiple	10.7×	**1.5×**	**0.6×**

> **The burn multiple story:** 10.7× → 1.5× → 0.6× — from pure product investment to best-in-class efficiency. This is the Series A pitch.

Path to Cash-Flow Positive

Total capital required: ~$13M ($3M seed + $10M Series A). Cash-flow positive by Q1 Year 4.

The Ask — $5–8M Seed

> **Updated from $3M** based on market maturity. Competitors are raising $50–130M rounds. A larger seed enables faster category capture and stronger Series A positioning.

Category	Amount ($5M)	Allocation

Engineering	$2.0M (40%)	5 engineers + CTO. Core platform, MCP security, anomaly detection ML
Go-To-Market	$0.8M (16%)	DevRel, OSS community, content. Enterprise AE starts Month 6
Infrastructure	$0.3M (6%)	Cloud, CI/CD, SOC 2 readiness
Operations	$0.4M (8%)	Legal, IP, finance
Working Capital	$0.5M (10%)	Enterprise AR buffer
Reserve	$1.0M (20%)	Buffer for extended runway / opportunistic hires
Total	**$5.0M**	**~20 months runway**

18-Month Milestones

Milestone	Target
OSS launch + 5 framework integrations	Month 4
MCP security layer shipped	Month 5
First enterprise design partner	Month 6
$500K pipeline	Month 10
25 design partners	Month 14
$1.5–2.5M ARR / Series A ready	Month 16–18

Compliance Framework Mapping

Framework	Key Requirements	AgentGuard Capability
EU AI Act	Risk mgmt (Art. 9), Records (Art. 12), Transparency (Art. 13), Human oversight (Art. 14)	Policy engine, audit trail, dashboards, kill switch
NIST AI RMF	GOVERN, MAP, MEASURE, MANAGE	Full lifecycle coverage
ISO 42001	Risk assessment, operational planning	Automated scoring, policy enforcement
DORA	ICT risk management (EU financial)	Resilience testing, incident reporting
SOC 2 Type II	Access, operations, change management	Controls, monitoring, audit trails
HIPAA	Technical safeguards	Data access policies, encryption
SEC Cyber Rules	Material incident disclosure	Incident determination, board reporting

Risks & Mitigations

Risk	Prob.	Mitigation
Market timing — adoption slower	30%	Lean burn; product also secures traditional LLM APIs
Platform bundling — "good enough" from AWS/Google	40%	Framework-agnostic + compliance depth they won't build
Enterprise sales cycles 6–9 months	50%	PLG revenue while enterprise gestates; departmental land
Pricing erosion from OSS alternatives	25%	Compliance modules = pricing moat; proprietary anomaly ML
Key person risk at 4–8 people	35%	4-year vesting, documentation, cross-training

The One-Liner

> *"Every company deploying AI agents will need AgentGuard — the same way every company with a network needed a firewall. We're building that firewall."*

Next Steps

1. **Validate** — 15 enterprise discovery calls with CISOs and AI platform leads
2. **Build MVP** — Open-source policy engine wrapping LangChain + OpenAI Assistants (8-week sprint)
3. **Prove** — Document 5 real agent failure modes as threat research + marketing content
4. **Fund** — Close $5M seed: angels + institutional lead
5. **Launch** — OSS release + Product Hunt + Hacker News + CISO roundtable series
6. **Partner** — LOIs from 3 enterprise design partners before seed close

AgentGuard — Business Case v2.0 — Confidential — February 2026
Includes CSO security review, CFO financial stress-test, and competitive market intelligence